




Milestone 2

Group JeLLY-B

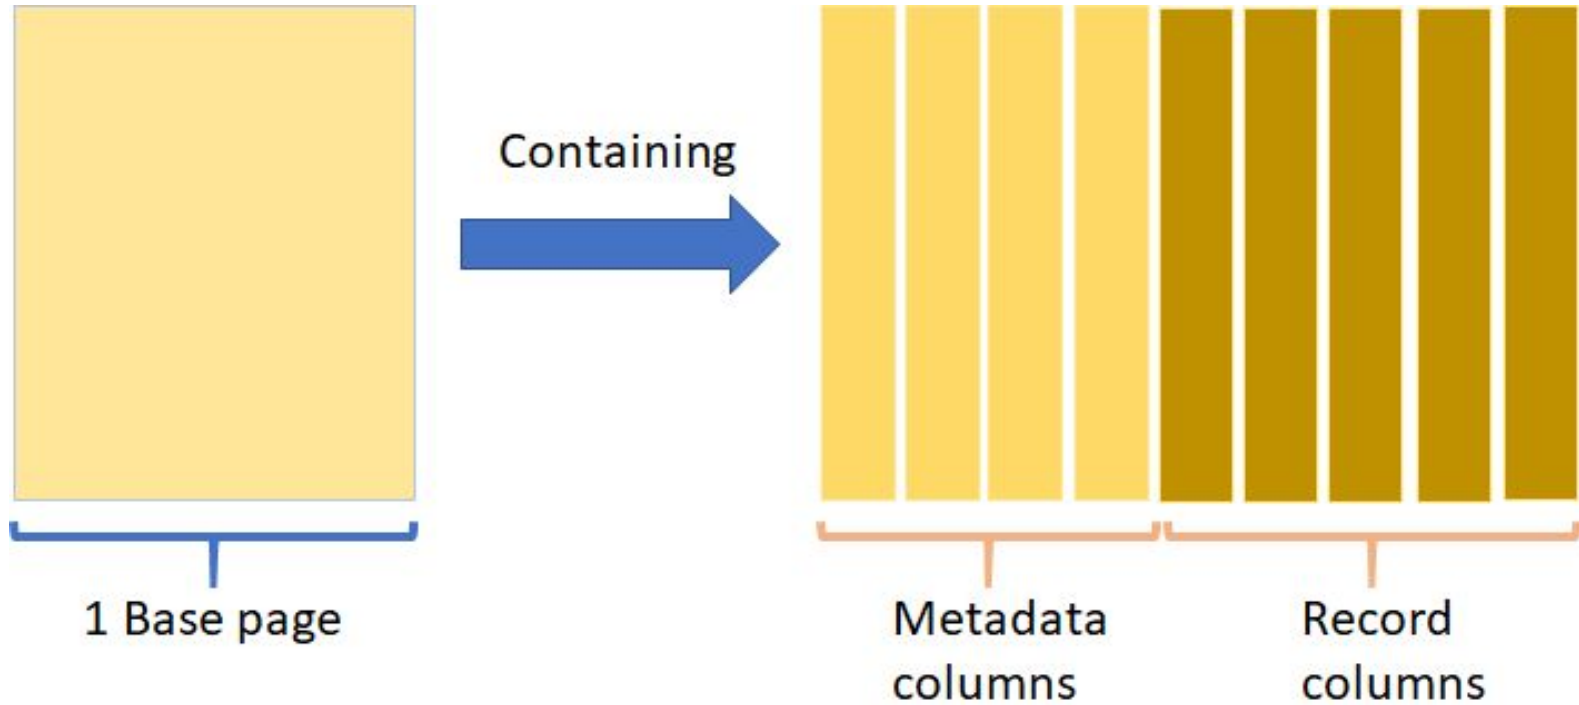


Jessica LaCourse
Lingxiao Meng
Lisa Malins
Yinuo Zhang
Benjamin Rausch

Page range



L-store structure



Durable Storage

All data are now pushed to the hard disk as needed

- Upon creation of a new database, a bufferpool is now initialized with the new table
- The data is pickled, or serialized, to improve space efficiency on disk
- Built-in functionality ensures enough space is allocated on disk for the new file before flushing from bufferpool to disk
- Files stored at page-level

Storage Structure

- Each page range stored in its own file
- Separation of concerns inspired by department stores: Take a Number!
 - Each “physical page” object is created with its table’s name and the index of its page range.
 - On creation, it asks the DB’s bufferpool for some space in its page range’s file.
 - The bufferpool returns the first offset in that file where a new page could be written.
- Physical page objects only store the “address” of its data on disk. It’s the bufferpool that accesses the data, going to disk if necessary.

Persistence

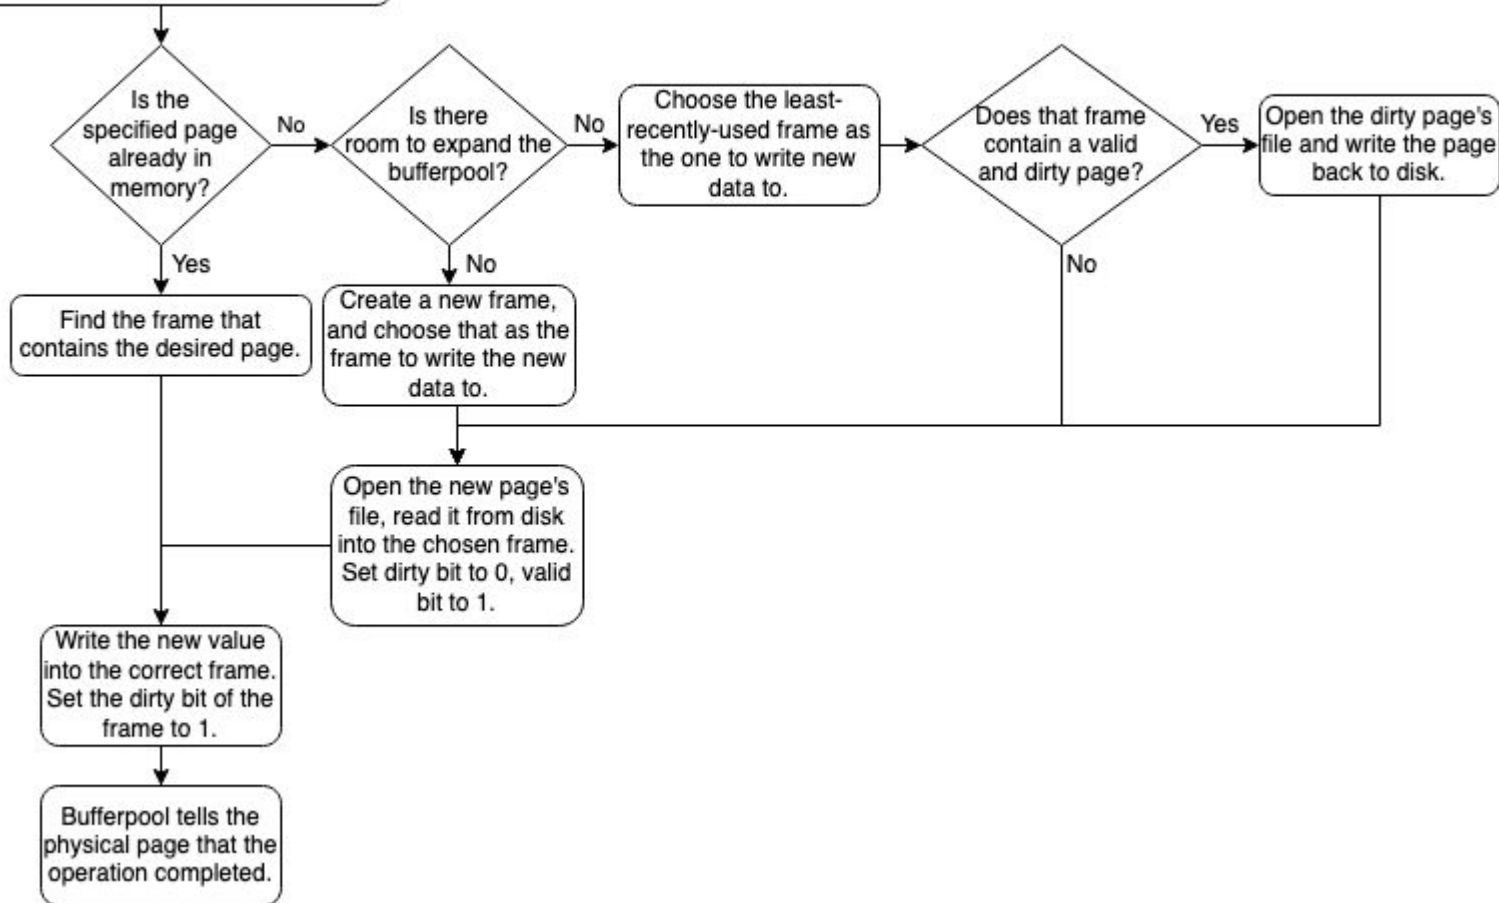
- If the database is closed, it destroys certain unnecessary structures, asks the bufferpool to persist dirty pages to disk, then pickles its members to a “backup” file on disk.
- Whenever the database is opened, it checks for this backup file.

Bufferpool

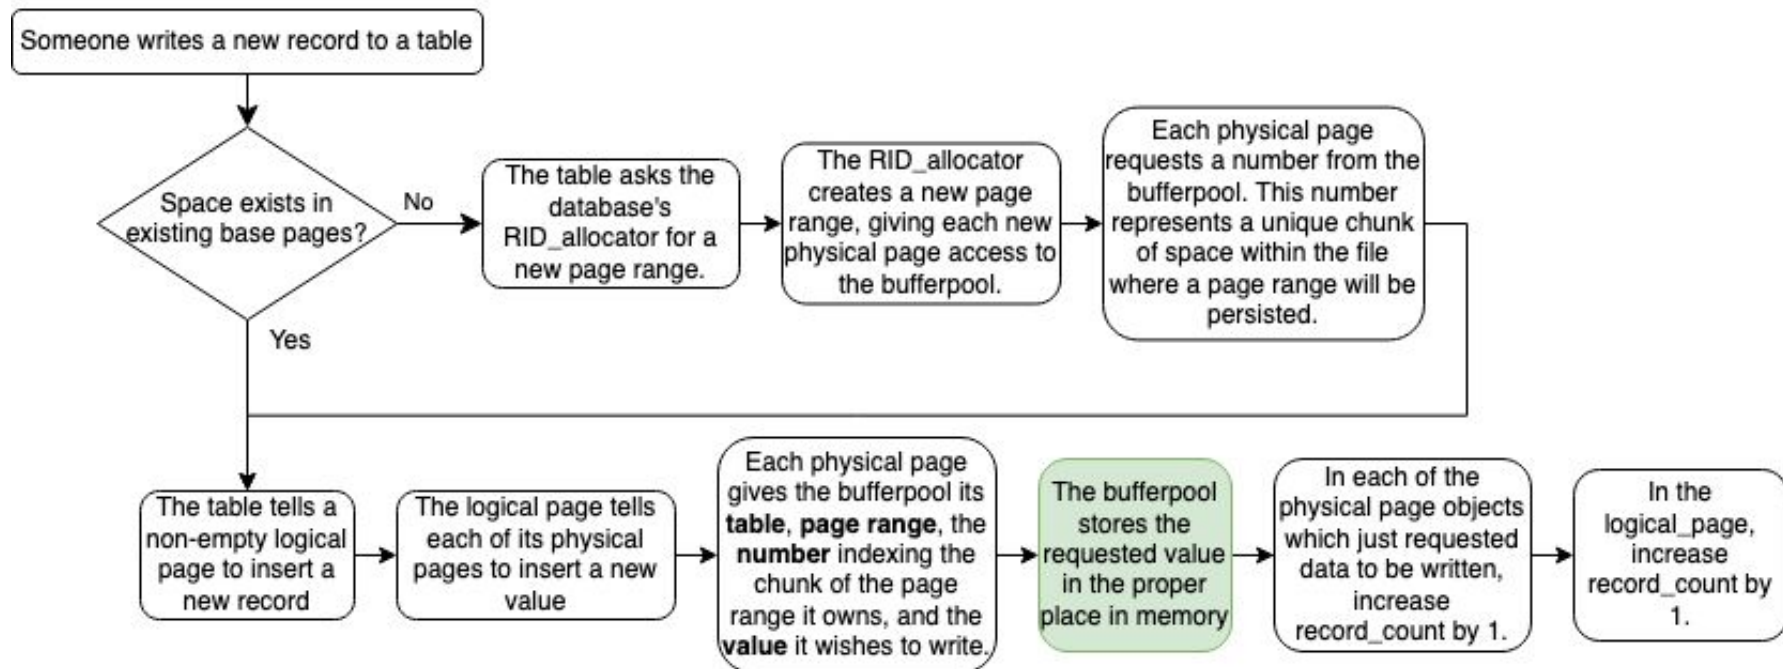
- Built to improve efficiency for read/write requests
 - Cache will be much quicker than referencing hard disk
- Each page brought into the bufferpool will be assigned a frame, indexed, and pinned until the respective query (or queries) no longer need to reference the page
- Uses a Least Recently Used (LRU) Cache Implementation to evict the oldest pages in order to make space for newly requested pages
- Pages to be evicted, dirty or not, are flushed to disk

Bufferpool Memory Management

A physical page needs to write a value.
It passes the request to the bufferpool,
identifying itself in the request.



Storage Example: Insertion



Pre-Merge

- Tail records have metadata column for base RID to aid in merge
- Within the insertion function:
Check_merge() stores page range ID's for full (complete) base pages
- Within the update function:
Once a tail page is full, its ID will be added to merge queue
- Merge queue will pick the first page that meets the requirements:
 - (1) Base pages within a given page range are full
 - (2) The tail page to be merged is attached to the same page range
 - (3) Otherwise tail pages are still stored in the queue (with FIFO order maintained)

Merge

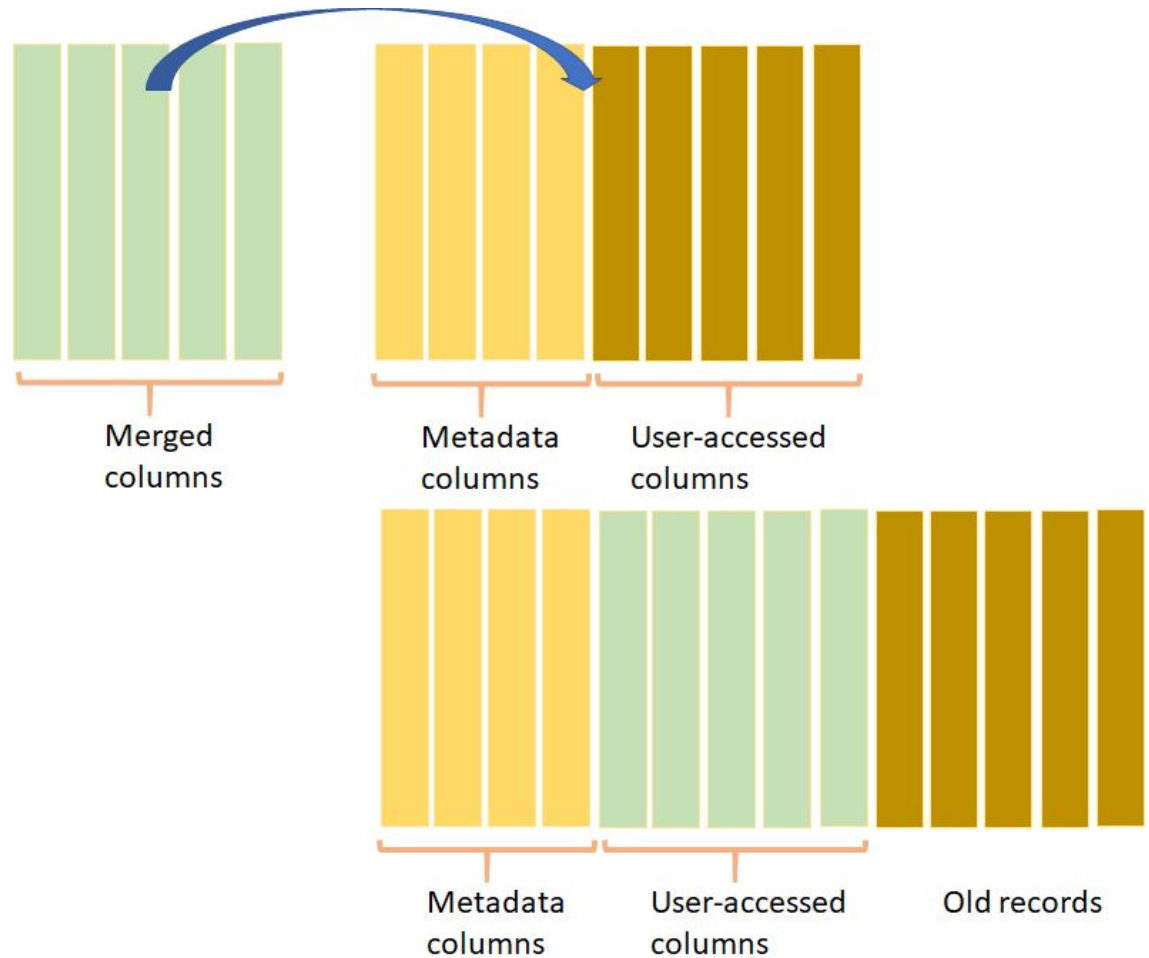
- Python threading functionality allows a merge thread in addition to the main thread
- These are core functions to ensure contention-free merges
 - (1) Ensures read-only records and eliminates risk of read-write conflicts
 - (2) Improves efficiency and accuracy by batching updates and pulling only the necessary pages for merge

Merge

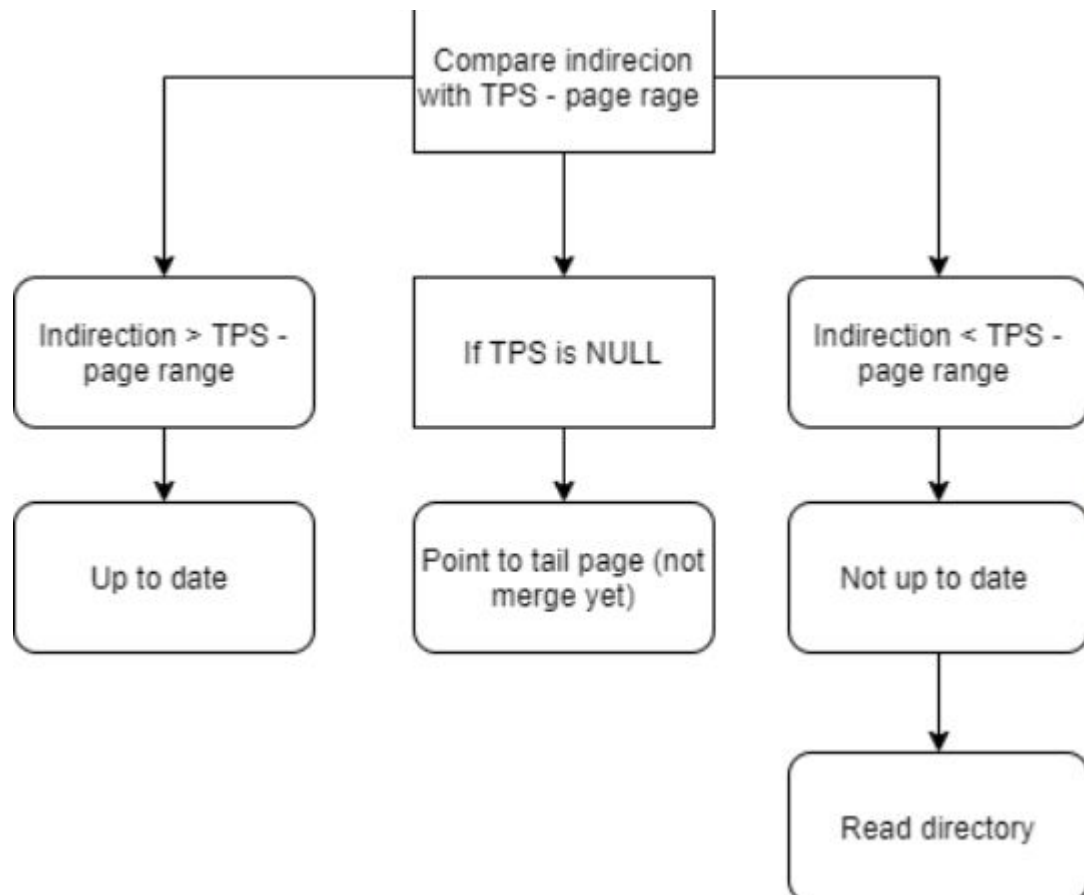
- Using cumulative updates, the latest tail entries will contain up-to-date information for their corresponding base records
- The original user data from base pages is copied
 - Metadata is left intact and untouched
 - Copy the most-up-to-date records if their base RID are marked updated
- Running concurrently with any user requests, merge() matches the updated records to their respective base record metadata
- The page directory is updated and a lock is placed while updating:
 - Users requesting information before and during the merge will receive older references
 - Users requesting information after the merge will receive up-to-date information

Page directory update

- Acquire lock during this process
- Releasing lock once finished
- Store TPS for merged page range
- (Original TPS is set to None)
- Old records are discarded



Selection



Indexing on any column

Milestone 1 implemented index on primary key column only

Now able to index on any column

- Index on primary key column: only 1 RID per primary key
- Index on non-primary key column: may have multiple RIDs per key
- $O(1)$ check to see if value exists in column
- Return all records with that value

Index

Insertion

- Let `list_of_RIDs_for_this_value = self.data[column].get(value)`
- If `list_of_RIDs_for_this_value` is none, get `self.data[column][value]`
- Update the insert data
- After this call, `self.locate` will return a list, which contains RID

Index

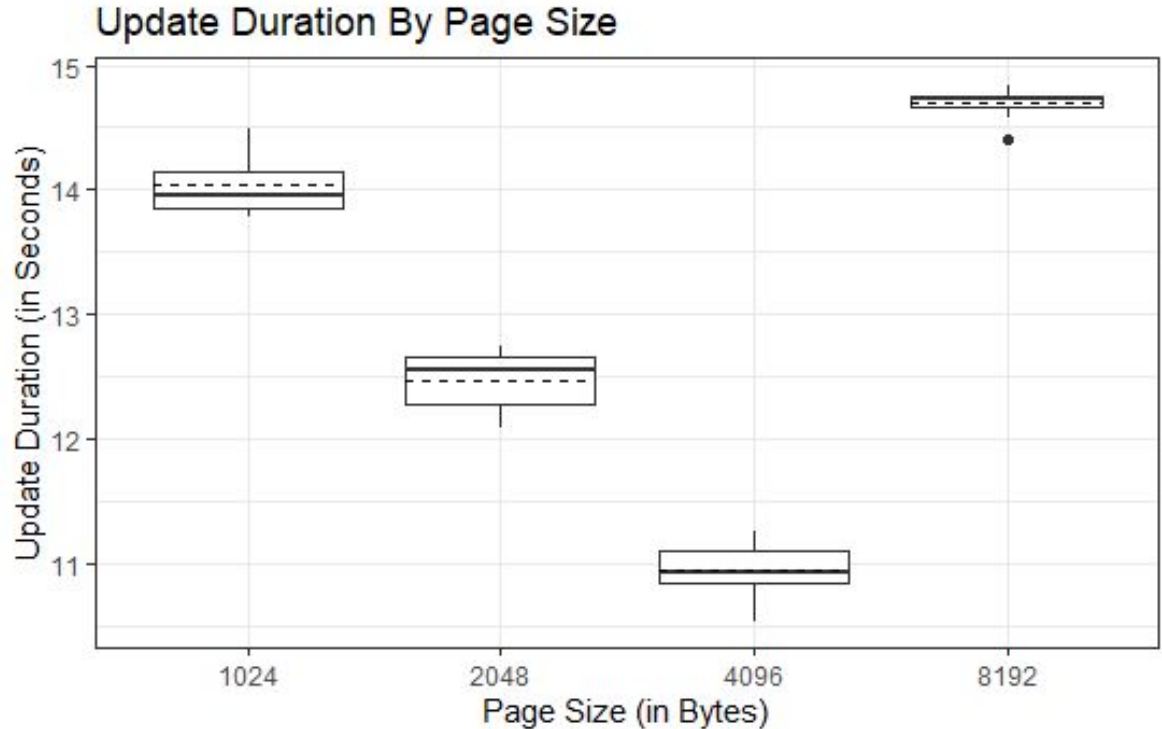
Deletion

- Use given value to get the RIDs
- Check the RIDs are associate with given value
- Remove these RIDs from index.
- After this call, `self.locate` will return a list, which does not contain RID

Performance

Optimal Page Size;
1,000 Records

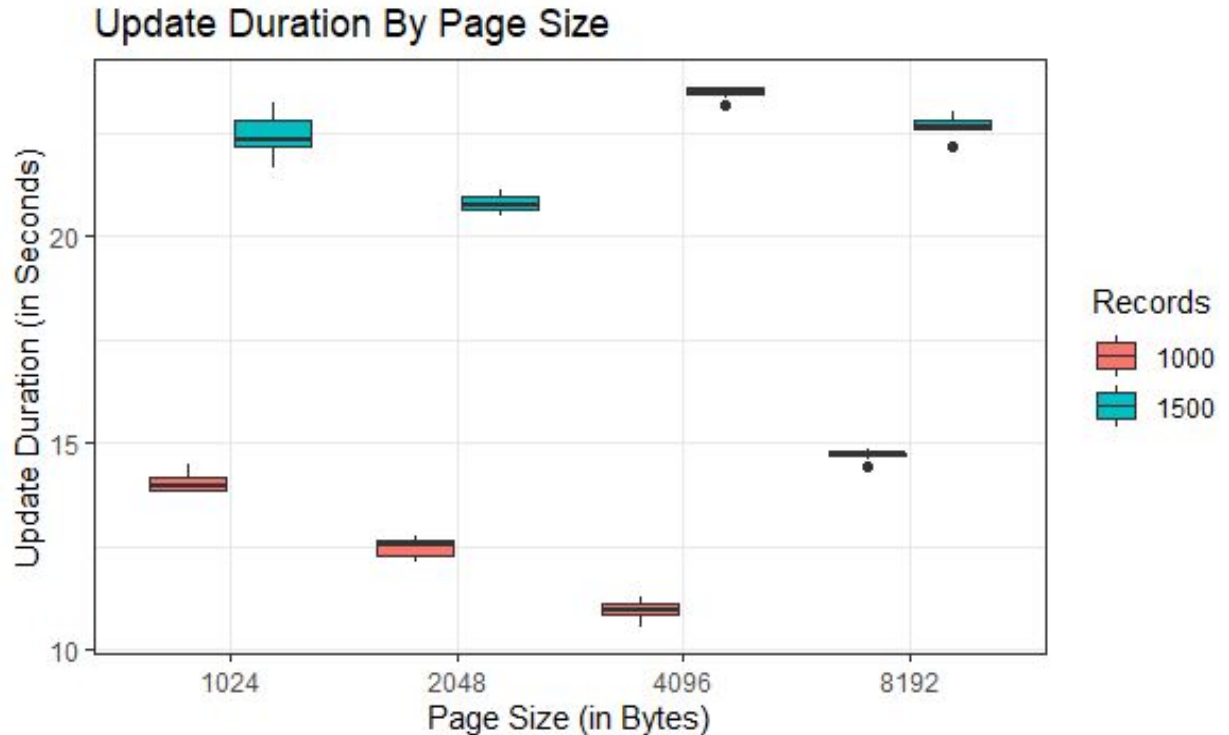
- Notable decrease in size up until 4kb/page
- Increase with 8kb/page showing impact when page(s) not entirely full



Performance

Optimal Page Size; Record Number Comparison

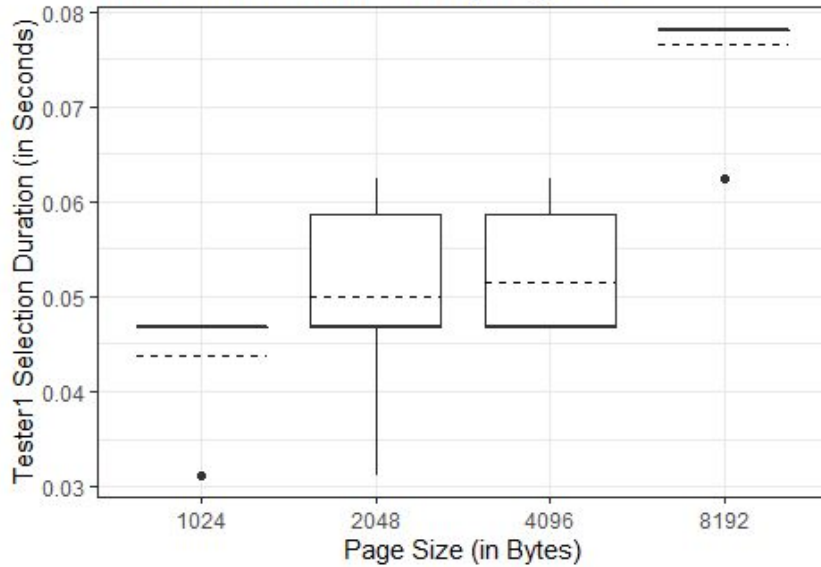
- Ratio between page size and number of records appears to directly affect performance
- Increases in the number of records (intuitively) increase processing time




Performance

Optimal Page Size; 1,000 Records

Tester 1 Selection Time By Page Size



Tester 2 Selection Time By Page Size



Processor Intel(R) Core(TM) i7-9750H CPU @ 2.60GHz 2.59 GHz

Installed RAM 16.0 GB